

Not Connected

Not Trusted

Python 2

File Edit View Insert Cell Kernel Widgets Help

Code

```
In [1]: 1 %reload_ext autoreload
2 %autoreload 2
3 %matplotlib inline
4 <
5 from fastai.io import *
6 from fastai.conv_learner import *
7 <
8 from fastai.column_data import *
```

Setup

We're going to download the collected works of Nietzsche to use as our data for this class.

```
In [2]: 1 PATH = 'data/nietzsche/'
```

```
In [3]: 1 get_data("https://s3.amazonaws.com/text-datasets/nietzsche.txt", f'{PATH}nietzsche.txt')
2 text = open(f'{PATH}nietzsche.txt').read()
3 print('corpus length:', len(text))
```

No of characters in the whole text

corpus length: 600893 show the first 400 characters of the text

```
In [4]: 1 text[:400]
```

```
Out[4]: 'PREFACE\n\n\nSUPPOSING that Truth is a woman--what then? Is there not gro
und\nfor suspecting that all philosophers, in so far as they have been\ndo
gmatisms, have failed to understand women--that the terrible\nseriousness
and clumsy importunity with which they have usually paid\ntheir addresses
to Truth, have been unskilled and unseemly methods for\nwinning a woman? C
ertainly she has never allowed herself '
```

```
In [5]: 1 chars = sorted(list(set(text)))
2 vocab_size = len(chars)+1
3 print('total chars:', vocab_size)
```

total chars: 85

set creates a list of unique characters, in this case 85 unique characters

Sometimes it's useful to have a zero value in the dataset, e.g. for padding

```
In [6]: 1 chars.insert(0, "\0")
```

insert \ padding

```
In [7]: 1 ''.join(chars[1:-6])
```

```
Out[7]: '\n !"\'(),-.0123456789:;=?ABCDEFGHIJKLMNOPQRSTUVWXYZ[_]_abcdefghijklmnopqrstuvwxyz
```

Map from chars to indices and back again

map every character to a unique ID and every unique ID to every character

```
In [8]: 1 char_indices = dict((c, i) for i, c in enumerate(chars))
2 indices_char = dict((i, c) for i, c in enumerate(chars))
```

idx will be the data we use from on the mapping above)

converts each character into a corresponding index

converts all the characters to their index (based

```
In [9]: 1 idx = [char_indices[c] for c in text]
```

collected works of Nietzsche

```
In [10]: 1 idx[:10]
```

```
Out[10]: [40, 42, 29, 30, 25, 27, 29, 1, 1, 1]
```

P R E F A C E n n n

```
In [11]: 1 ''.join(indices_char[i] for i in idx[:70])
```

```
Out[11]: 'PREFACE\n\n\nSUPPOSING that Truth is a woman--what then? Is there not gro
```

Three char model

Create inputs

Create a list of every 4th character, starting at the 0th, 1st, 2nd, then 3rd characters

0th character

1st character

2nd character

3rd character

```
1 cs=3
2 c1_dat = [idx[i] for i in range(0, len(idx)-1-cs, cs)]
3 c2_dat = [idx[i+1] for i in range(0, len(idx)-1-cs, cs)]
4 c3_dat = [idx[i+2] for i in range(0, len(idx)-1-cs, cs)]
5 c4_dat = [idx[i+3] for i in range(0, len(idx)-1-cs, cs)]
```

skip every 3 characters

Our inputs

```
In [13]: 1 x1 = np.stack(c1_dat[:-2])
2 x2 = np.stack(c2_dat[:-2])
3 x3 = np.stack(c3_dat[:-2])
```

0th, 1st and 2nd character as inputs

Our output

```
In [14]: 1 y = np.stack(c4_dat[:-2])
```

predict the 3rd character as the output

The first 4 inputs and outputs

```
In [14]: 1 x1[:4], x2[:4], x3[:4]
```

```
Out[14]: (array([40, 30, 29, 1]), array([42, 25, 1, 43]), array([29, 27, 1, 45])
          )
```

0th P 1st R 2nd E

```
In [15]: 1 y[:4]
```

```
Out[15]: array([30, 29, 1, 40])
```

F

used to predict the 3rd character

```
In [16]: 1 x1.shape, y.shape
```

```
Out[16]: ((200295,), (200295,))
```

Create and train model

Pick a size for our hidden state

```
In [15]: 1 n_hidden = 256
```

The number of latent factors to create (i.e. the size of the embedding matrix)

```
In [16]: 1 n_fac = 42
```

1/2 the number of characters

```
In [30]: 1 class Char3Model(nn.Module):
2         def __init__(self, vocab_size, n_fac):
3             super().__init__()
4             self.e = nn.Embedding(vocab_size, n_fac)
5         <
6             # The 'green arrow' from our diagram - the layer operation from
             hidden
7             self.l_in = nn.Linear(n_fac, n_hidden)
8         <
9             # The 'orange arrow' from our diagram - the layer operation from
             to hidden
10            self.l_hidden = nn.Linear(n_hidden, n_hidden)
11
12            # The 'blue arrow' from our diagram - the layer operation from i
             output
13            self.l_out = nn.Linear(n_hidden, vocab_size)
14
15            def forward(self, c1, c2, c3):
16                in1 = F.relu(self.l_in(self.e(c1)))
17                in2 = F.relu(self.l_in(self.e(c2)))
18                in3 = F.relu(self.l_in(self.e(c3)))
19
20                h = V(torch.zeros(in1.size()).cuda())
21                h = F.tanh(self.l_hidden(h+in1))
22                h = F.tanh(self.l_hidden(h+in2))
23                h = F.tanh(self.l_hidden(h+in3))
24
25                return F.log_softmax(self.l_out(h))
```

square matrix

when calling forward we will pass 3 characters c1, c2, and c3

each character will pass through an embedding 'e', linear layer 'l' and a relu

char 3

char 2

char 1

lin3

lin2

lin1

```
In [20]: 1 md = ColumnarModelData.from_arrays('.', [-1], np.stack([x1,x2,x3], axis=
```

```
, bs=512)
```

hyperbolic tanh

```
In [21]: 1 m = Char3Model(vocab_size, n_fac).cuda()
```

```
In [22]: 1 it = iter(md.trn_dl)
2 *xs, yt = next(it)
3 t = m(*V(xs))
```

inputs of 3 and size 512

Variablized versions of the inputs

```
In [191]: 1 opt = optim.Adam(m.parameters(), 1e-2)
```

```
In [192]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          2.09627  6.52849]
```

```
In [193]: 1 set_lrs(opt, 0.001)
```

```
In [194]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          1.84525  6.52312]
```

Test model

```
In [195]: 1 def get_next(inp):
2         idxs = T(np.array([char_indices[c] for c in inp]))
3         p = m(*VV(idxs))
4         i = np.argmax(to_np(p))
5         return chars[i]
```

convert to tensor T

```
In [196]: 1 get_next('y. ')
```

3 char model hence can pass in 3 things

```
Out[196]: 'T'
```

```
In [197]: 1 get_next('ppl')
```

```
Out[197]: 'e'
```

```
In [198]: 1 get_next(' th')
```

```
Out[198]: 'e'
```


```
In [199]: 1 get_next('and')
```

```
Out[199]: ' '
```

Our first RNN!

Create inputs

8 layer network =
8 time steps in
RNN



This is the size of our unrolled RNN

we are now going to
predict every 8th
character

```
In [19]: 1 cs=8
```

For each of 0 through 7, create a list of every 8th character with that starting point. These will be the 8 inputs to our model.

```
In [20]: 1 c_in_dat = [[idx[i+j] for i in range(cs) for j in range(len(idx)-cs-1)]]
```

Then create a list of the next character in each of these series. This will be the labels for our model.

```
In [21]: 1 c_out_dat = [idx[j+cs] for j in range(len(idx)-cs-1)]
```

```
In [22]: 1 xs = np.stack(c_in_dat, axis=0)
```

```
In [23]: 1 xs.shape
```

corpus length

```
Out[23]: (600884, 8)
```

char size

```
In [24]: 1 y = np.stack(c_out_dat)
```

So each column below is one series of 8 characters from the text.

input

```
In [25]: 1 xs[:,cs,:cs]
```

```
Out[25]: array([[40, 42, 29, 30, 25, 27, 29, 1],
 [42, 29, 30, 25, 27, 29, 1, 1],
 [29, 30, 25, 27, 29, 1, 1, 1],
 [30, 25, 27, 29, 1, 1, 1, 43],
 [25, 27, 29, 1, 1, 1, 43, 45],
 [27, 29, 1, 1, 1, 43, 45, 40],
 [29, 1, 1, 1, 43, 45, 40, 40],
 [1, 1, 1, 43, 45, 40, 40, 39]])
```

characters 0 to 7

characters 1 to 8 etc

output

...and this is the next character after each sequence.

```
In [26]: 1 y[:,cs]
```

```
Out[26]: array([ 1, 1, 43, 45, 40, 40, 39, 43])
```

the y prediction is the
8th character

Create and train model

```
In [27]: 1 val_idx = get_cv_idxes(len(idx)-cs-1)
```

```
In [28]: 1 md = ColumnarModelData.from_arrays('.', val_idx, xs, y, bs=512)
```

```
In [158]: 1 class CharLoopModel(nn.Module):
2     def __init__(self, vocab_size, n_fac):
3         super().__init__()
4         self.e = nn.Embedding(vocab_size, n_fac)
5         self.l_in = nn.Linear(n_fac, n_hidden)
6         self.l_hidden = nn.Linear(n_hidden, n_hidden)
7         self.l_out = nn.Linear(n_hidden, vocab_size)
8
9     def forward(self, *cs):
10        bs = cs[0].size(0)
11        h = V(torch.zeros(bs, n_hidden).cuda())
12        for c in cs:
13            inp = F.relu(self.l_in(self.e(c)))
14            h = F.tanh(self.l_hidden(h+inp))
15
16        return F.log_softmax(self.l_out(h))
```

same as 3 char model
above except now in a
loop or RNN

```
In [159]: 1 m = CharLoopModel(vocab_size, n_fac).cuda()
2         opt = optim.Adam(m.parameters(), 1e-2)
```

```
In [160]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

[0. 2.01881 2.0294]

```
In [161]: 1 set_lrs(opt, 0.001)
```

```
In [162]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

[0. 1.7161 1.7246]

loss

```
In [163]: 1 class CharLoopConcatModel(nn.Module):
2     def __init__(self, vocab_size, n_fac):
3         super().__init__()
4         self.e = nn.Embedding(vocab_size, n_fac)
5         self.l_in = nn.Linear(n_fac+n_hidden, n_hidden)
6         self.l_hidden = nn.Linear(n_hidden, n_hidden)
7         self.l_out = nn.Linear(n_hidden, vocab_size)
8
9     def forward(self, *cs):
10        bs = cs[0].size(0)
11        h = V(torch.zeros(bs, n_hidden).cuda())
12        for c in cs:
13            inp = torch.cat((h, self.e(c)), 1)
14            inp = F.relu(self.l_in(inp))
15            h = F.tanh(self.l_hidden(inp))
16
17        return F.log_softmax(self.l_out(h))
```

concatenating is a
good design heuristic if
you have different
types of information
you want to combine.

Adding information
even if same shape
generally means losing
information

h = hidden layer here
is a rank 2 tensor = 2
dimensions

starting point for 'for' loop

because concatenating
now linear layer is
n_fac+n_hidden to
n_hidden instead of
n_fac to n_hidden as
before

this is now size n_fac +
n_hidden

makes it back to size
n_hidden

same square matrix as
before so still size
n_hidden

```
In [164]: 1 m = CharLoopConcatModel(vocab_size, n_fac).cuda()  
2 opt = optim.Adam(m.parameters(), 1e-3)
```

```
In [165]: 1 it = iter(md.trn_dl)  
2 *xs, yt = next(it)  
3 t = m(*V(xs))
```

```
In [166]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          1.80699  1.7688 ]
```

```
In [167]: 1 set_lrs(opt, 1e-4)
```

```
In [168]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          1.69154  1.68602 ]
```

concatenating reduces
loss

Test model

```
In [46]: 1 def get_next(inp):  
2     idxs = T(np.array([char_indices[c] for c in inp]))  
3     p = m(*VV(idxs))  
4     i = np.argmax(to_np(p))  
5     return chars[i]
```

```
In [47]: 1 get_next('for thos')
```

can now pass in 8
chars

```
Out[47]: 'e'
```

```
In [48]: 1 get_next('part of ')
```

```
Out[48]: 't'
```

```
In [49]: 1 get_next('queens a')
```

```
Out[49]: 'n'
```

RNN with pytorch

same as before but
using pytorch and less
code

```
In [169]: 1 class CharRnn(nn.Module):  
2     def __init__(self, vocab_size, n_fac):  
3         super().__init__()  
4         self.e = nn.Embedding(vocab_size, n_fac)  
5         self.rnn = nn.RNN(n_fac, n_hidden)  
6         self.l_out = nn.Linear(n_hidden, vocab_size)  
7
```

```

8     def forward(self, *cs):
9         bs = cs[0].size(0)
10        h = V(torch.zeros(1, bs, n_hidden))
11        inp = self.e(torch.stack(cs))
12        outp, h = self.rnn(inp, h)
13
14        return F.log_softmax(self.l_out(outp[-1]))

```

h - hidden layer here is a rank 3 tensor = 3 dimensions

starting point for 'for' loop

```

In [170]: 1 m = CharRnn(vocab_size, n_fac).cuda()
          2 opt = optim.Adam(m.parameters(), 1e-3)

```

```

In [171]: 1 it = iter(md.trn_dl)
          2 *xs, yt = next(it)

```

```

In [172]: 1 t = m.e(V(torch.stack(xs)))
          2 t.size()

```

Out[172]: torch.Size([8, 512, 42])

char size

batch size

size of the embedding matrix n_fac

```

In [173]: 1 ht = V(torch.zeros(1, 512, n_hidden))
          2 outp, hn = m.rnn(t, ht)
          3 outp.size(), hn.size()

```

Out[173]: (torch.Size([8, 512, 256]), torch.Size([1, 512, 256]))

unit access = allows RNNs to be bi-directional in this case it is a 1 hence the reason the output is 1

```

In [175]: 1 t = m(*V(xs)); t.size()

```

n_hidden

Out[175]: torch.Size([512, 85])

```

In [176]: 1 fit(m, md, 4, opt, F.nll_loss)

```

A Jupyter Widget

```

[ 0.      1.86162  1.84379]
[ 1.      1.66857  1.66926]
[ 2.      1.5879   1.60001]
[ 3.      1.55005  1.55722]

```

```

In [177]: 1 set_lrs(opt, 1e-4)

```

```

In [178]: 1 fit(m, md, 2, opt, F.nll_loss)

```

A Jupyter Widget

```

[ 0.      1.45798  1.5094 ]
[ 1.      1.45333  1.50419]

```

loss improves

Test model

```

In [179]: 1 def get_next(inp):
          2     idxs = T(np.array([char_indices[c] for c in inp]))
          3     p = m(*VV(idxs))

```



```
4 i = np.argmax(to_np(p))
5 return chars[i]
```

```
In [180]: 1 get_next('for thos')
```

```
Out[180]: 'e'
```

```
In [181]: 1 def get_next_n(inp, n):
2         res = inp
3         for i in range(n):
4             c = get_next(inp)
5             res += c
6             inp = inp[1:]+c
7         return res
```

```
In [182]: 1 get_next_n('for thos', 40)
```

```
Out[182]: 'for those of the same to the same to the same to'
```

Multi-output model

Setup

Let's take non-overlapping sets of characters this time

```
In [29]: 1 c_in_dat = [[idx[i+j] for i in range(cs)] for j in range(0, len(idx)-cs
-1, cs)]
```

Then create the exact same thing, offset by 1, as our labels

```
In [30]: 1 c_out_dat = [[idx[i+j] for i in range(cs)] for j in range(1, len(idx)-
cs, cs)]
```

```
In [31]: 1 xs = np.stack(c_in_dat)
2 xs.shape
```

```
Out[31]: (75111, 8)
```

shape now 75111 long
instead of 600884 as
we are looking every 8
non-overlapping
characters (600884/8 =
75111)

```
In [32]: 1 ys = np.stack(c_out_dat)
2 ys.shape
```

```
Out[32]: (75111, 8)
```

```
In [33]: 1 xs[:cs,:cs]
```

In this case using non-
overlapping characters

```
Out[33]: array([[40, 42, 29, 30, 25, 27, 29, 1],
 [ 1,  1, 43, 45, 40, 40, 39, 43],
 [33, 38, 31,  2, 73, 61, 54, 73],
 [ 2, 44, 71, 74, 73, 61,  2, 62],
 [72,  2, 54,  2, 76, 68, 66, 54],
 [67,  9,  9, 76, 61, 54, 73,  2],
```

first 8 characters

next 8 characters

```
[73, 61, 58, 67, 24, 2, 33, 72],
[ 2, 73, 61, 58, 71, 58, 2, 67]])
```

```
In [34]: 1 ys[:cs,:cs]
```

```
Out[34]: array([[42, 29, 30, 25, 27, 29, 1, 1],
 [ 1, 43, 45, 40, 40, 39, 43, 33],
 [38, 31, 2, 73, 61, 54, 73, 2],
 [44, 71, 74, 73, 61, 2, 62, 72],
 [ 2, 54, 2, 76, 68, 66, 54, 67],
 [ 9, 9, 76, 61, 54, 73, 2, 73],
 [61, 58, 67, 24, 2, 33, 72, 2],
 [73, 61, 58, 71, 58, 2, 67, 68]])
```

Create and train model

```
In [35]: 1 val_idx = get_cv_idxxs(len(xs)-cs-1)
```

```
In [36]: 1 md = ColumnarModelData.from_arrays('.', val_idx, xs, ys, bs=512)
```

```
In [136]: 1 class CharSeqRnn(nn.Module):
2     def __init__(self, vocab_size, n_fac):
3         super().__init__()
4         self.e = nn.Embedding(vocab_size, n_fac)
5         self.rnn = nn.RNN(n_fac, n_hidden)
6         self.l_out = nn.Linear(n_hidden, vocab_size)
7
8     def forward(self, *cs):
9         bs = cs[0].size(0)
10        h = V(torch.zeros(1, bs, n_hidden))
11        inp = self.e(torch.stack(cs))
12        outp,h = self.rnn(inp, h)
13        return F.log_softmax(self.l_out(outp))
```

```
In [137]: 1 m = CharSeqRnn(vocab_size, n_fac).cuda()
2 opt = optim.Adam(m.parameters(), 1e-3)
```

2nd axis =
batch size
= 512

```
1 it = iter(md.trn_dl)
2 *xst,yt = next(it)
```

have to write a custom
loss function because
pytorch takes into
account tensor ranks

```
In [139]: 1 def nll_loss_seq(inp, targ):
2     sl,bs,nh = inp.size()
3     targ = targ.transpose(0,1).contiguous().view(-1)
4     return F.nll_loss(inp.view(-1,nh), targ)
```

1st axis = sl =
sequence length of
RNN in this case 8
characters = time steps

3rd axis = hidden state
= 256

flatten targets

flatten inputs

```
fit(m, md, 4, opt, nll_loss_seq)
```

A Jupyter Widget

```
[ 0.      1.1164  0.9535]
[ 1.      0.83636 0.76285]
[ 2.      0.71097 0.68187]
[ 3.      0.64594 0.63435]
```

```
In [141]: 1 set_lrs(opt, 1e-4)
```


```
In [142]: 1 fit(m, md, 1, opt, nll_loss_seq)
```

A Jupyter Widget

```
[ 0.      0
```

Identity init!

example size 3 identity matrix trick about identity matrixes is that something times an identity matrix produces itself preventing gradient explosion



```
In [145]: 1 m = CharSeqRnn(vocab_size, n_fac).cuda()  
2 opt = optim.Adam(m.parameters(), 1e-2)
```

```
In [146]: 1 m.rnn.weight_hh_l0.data.copy_(torch.eye(n_hidden))
```

eye for identity

Out [146]:

```
1  0  0 ... 0 0 0  
0  1  0 ... 0 0 0  
0  0  1 ... 0 0 0  
...  
0  0  0 ... 1 0 0  
0  0  0 ... 0 1 0  
0  0  0 ... 0 0 1
```

identity matrix of size n-hidden

use ?m.rnn to get info on a matrix; inputs, outputs, attributes

[torch.cuda.FloatTensor of size 256x256 (GPU 0)]

```
In [147]: 1 fit(m, md, 4, opt, nll_loss_seq)
```

A Jupyter Widget

```
[ 0.      0.80381  0.7284 ]  
[ 1.      0.66663  0.63739]  
[ 2.      0.61271  0.6188 ]  
[ 3.      0.59271  0.60541]
```

```
In [148]: 1 set_lrs(opt, 1e-3)
```

```
In [149]: 1 fit(m, md, 4, opt, nll_loss_seq)
```

A Jupyter Widget

```
[ 0.      0.5156  0.53407]  
[ 1.      0.49554  0.52318]  
[ 2.      0.48124  0.51721]  
[ 3.      0.47085  0.51224]
```

```
In [150]: 1 set_lrs(opt, 1e-4)
```

```
In [151]: 1 fit(m, md, 4, opt, nll_loss_seq)
```

A Jupyter Widget

```
[ 0.      0.4549  0.50927]
[ 1.      0.45277  0.50849]
[ 2.      0.45234  0.50804]
[ 3.      0.45165  0.50758]
```

Stateful model

Setup

```
In [156]: 1 from torchtext import vocab, data
          2 â€‹
          3 from fastai.nlp import *
          4 from fastai.lm_rnn import *
```

```
In [157]: 1 PATH='data/nietzsche/'
          2 â€‹
          3 TRN_PATH = 'trn/'
          4 VAL_PATH = 'val/'
          5 TRN = f'{PATH}{TRN_PATH}'
          6 VAL = f'{PATH}{VAL_PATH}'
          7 â€‹
          8 %ls {PATH}
```

```
nietzsche.txt  trn/  val/
```

```
In [158]: 1 TEXT = data.Field(lower=True, tokenize=list)
```

```
In [159]: 1 bs=64; bptt=8; n_fac=42; n_hidden=256
```

```
In [160]: 1 FILES = dict(train=TRN_PATH, validation=VAL_PATH, test=VAL_PATH)
          2 md = LanguageModelData.from_text_files(PATH, TEXT, **FILES, bs=bs, bptt=
          n_freq=3)
```

```
In [161]: 1 len(md.trn_dl), md.nt, len(md.trn_ds), len(md.trn_ds[0].text)
```

```
Out[161]: (963, 56, 1, 493747)
```

```
In [162]: 1 class CharSeqStatefulRnn(nn.Module):
          2     def __init__(self, vocab_size, n_fac, bs):
          3         super().__init__()
          4         self.e = nn.Embedding(vocab_size, n_fac)
          5         self.rnn = nn.RNN(n_fac, n_hidden)
          6         self.l_out = nn.Linear(n_hidden, vocab_size)
          7         self.h = V(torch.zeros(1, bs, n_hidden))
          8
          9     def forward(self, cs):
         10         bs = cs[0].size(0)
         11         if self.h.size(1) != bs: self.h = V(torch.zeros(1, bs, n_hidden))
         12         inp = self.e(cs)
         13         outp,h = self.rnn(inp, self.h)
```

```
14         self.h = repackage_var(h)
15         return F.log_softmax(self.l_out(outp)).view(-1,vocab_size)
```

```
In [172]: 1 m = CharSeqStatefulRnn(vocab_size, n_fac, 512).cuda()
          2 opt = optim.Adam(m.parameters(), 1e-3)
```

```
In [173]: 1 m.rnn.weight_hh_l0.data.copy_(torch.eye(n_hidden));
```

```
In [167]: 1 fit(m, md, 1, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          1.13513  1.08475]
```

```
In [174]: 1 set_lrs(opt, 1e-3)
```

```
In [175]: 1 fit(m, md, 2, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          0.96742  0.96103]
[ 1.          0.88555  0.8695 ]
```

```
In [176]: 1 fit(m, md, 2, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          0.87285  0.881   ]
[ 1.          0.87575  0.76261]
```

```
In [177]: 1 fit(m, md, 4, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          0.80812  0.82194]
[ 1.          0.85778  0.80804]
[ 2.          0.79613  0.73514]
[ 3.          0.827    0.78421]
```

```
In [178]: 1 fit(m, md, 4, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.          0.73793  0.76126]
[ 1.          0.74701  0.7797 ]
[ 2.          0.69897  0.75077]
[ 3.          0.69903  0.73123]
```

```
In [179]: 1 set_lrs(opt, 1e-4)
```

```
In [180]: 1 fit(m, md, 4, opt, F.nll_loss)
```

A Jupyter Widget

```
[ 0.      0.65736  0.72496]
[ 1.      0.66048  0.74494]
[ 2.      0.68257  0.75462]
[ 3.      0.69551  0.76302]
```

Test

```
In [181]: 1 def get_next(inp):
          2     idxs = TEXT.numericalize(inp)
          3     p = m(VV(idxs.transpose(0,1)))
          4     #     print(p)
          5     r = np.argmax(to_np(p), axis=1)
          6     #     print(r.shape)
          7     return TEXT.vocab.itos[r[-1]]
```

```
In [182]: 1 get_next('for thos')
```

Out[182]: 'k'

```
In [181]: 1 def get_next_n(inp, n):
          2     res = inp
          3     for i in range(n):
          4         c = get_next(inp)
          5         res += c
          6         inp = inp[1:]+c
          7     return res
```

```
In [182]: 1 get_next_n('for thos', 40)
```

Out[182]: 'for those of the same to the same to the same to'

```
In [ ]: 1 â€œ
```